

2. Data Manipulation

These data manipulation techniques are crucial for data preparation and analysis in various domains, from business analytics to data science and machine learning. They help ensure data quality, consistency, and usability for decision-making and reporting.

2.1 Import, store, and export data:

Fundamental understanding of ETL (Extract, Transform, Load): ETL is a process used in data management to extract data from various sources, transform it into a consistent format, and load it into a destination such as a data warehouse or a database. It involves the following steps:

- **Extract:** This step involves extracting data from various sources like databases, flat files, APIs, or web scraping.

example:

```
df.to_csv("D:/MyEduSolve/tugas_cleansing.csv")
```

extract the file into .csv format

```
df.to_csv("D:/MyEduSolve/tugas_cleansing.csv")
```

- **Transform:** In the transformation step, data is cleaned, standardized, and converted into a format that can be used for analysis. This might include data cleansing, data enrichment, and the creation of new derived variables.

example :

```
df['Referral'] = df['Referral'].astype(str)
df['Referral'] = df['Referral'].replace({'1.0': 'use Referral code', '0.0': 'not use Referral code'})
```

- **Load:** The final step is to load the transformed data into a target storage system, like a database or a data warehouse, making it available for analysis.

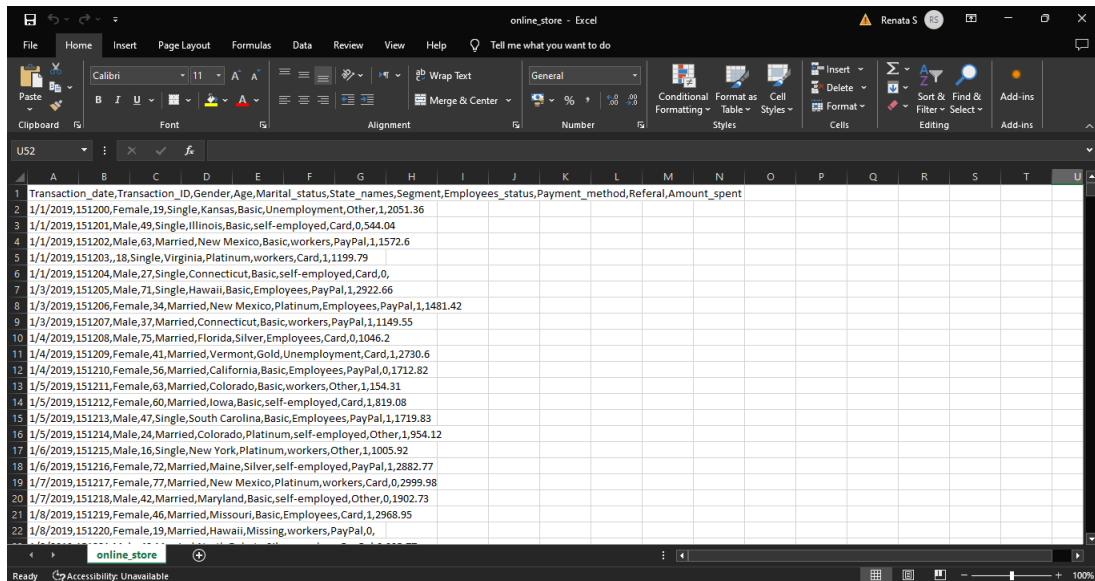
Load a CSV file containing datasets from online stores obtained from Kaggle

```
In [3]: df = pd.read_csv('online_store.csv')
```

- **Common data storage file formats:** These include delimited data files (e.g., CSV), XML (Extensible Markup Language), and JSON (JavaScript Object Notation) for storing structured data.

Here we use a CSV format file (Comma Separated Values), which means that each value in the data row in the file is separated by commas. Files in this format are generally used to store datasets.

```
In [3]: df = pd.read_csv('online_store.csv')
```



2.2 Clean data:

Purpose and common practices:

- **Handling NULL values:** Dealing with missing or NULL values, which may involve imputing missing data or excluding rows with missing values.

example:


Check data condition

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2512 entries, 0 to 2511
Data columns (total 11 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Transaction_date     2512 non-null   object
1   Transaction_ID       2512 non-null   int64
2   Gender               2484 non-null   object
3   Age                  2470 non-null   float64
4   Marital_status       2512 non-null   object
5   State_names          2512 non-null   object
6   Segment              2512 non-null   object
7   Employees_status     2486 non-null   object
8   Payment_method       2512 non-null   object
9   Referral             2357 non-null   float64
10  Amount_spent         2270 non-null   float64
dtypes: float64(3), int64(1), object(7)
memory usage: 216.0+ KB
```

Display a visualization of the columns Age.

`df.Age.plot(kind='hist')`



Because the Age column has a skewness distribution

Then we will do imputation on the Age column using the median

`val = df.Age.median()`

`df['Age'] = df.Age.fillna(val)`

Display dataset info to see whether the Age column has been imputed

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2512 entries, 0 to 2511
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Transaction_date       2512 non-null   object
1   Transaction_ID         2512 non-null   int64
2   Gender                 2484 non-null   object
3   Age                    2512 non-null   float64
4   Marital_status        2512 non-null   object
5   State_names            2512 non-null   object
6   Segment                2512 non-null   object
7   Employees_status      2486 non-null   object
8   Payment_method        2512 non-null   object
9   Referral               2357 non-null   float64
10  Amount_spent           2270 non-null   float64
dtypes: float64(3), int64(1), object(7)
memory usage: 216.0+ KB
```

From the dataset info above, it can be seen that the Age column has changed

- **Special characters:** Removing or encoding special characters that can cause data processing issues.

example :

You have the file name

```
df = pd.read_csv('online_store!!.csv')
```

And the double exclamation mark special character (!!) can cause problems when trying to read or process such files. You can remove these special character to make the file name cleaner and more easily accessible.

After removing special characters, the filename will become

```
df = pd.read_csv('online_store.csv')
```

- **Trimming spaces:** Trimming leading and trailing white spaces from text data to ensure consistency.

example:

Suppose you have text “ jupyter notebook “ that has extra spaces at the front and at the back. By trimming the extra spaces, the text will become:

“jupyter notebook”

This ensures that the text does not have unnecessary spaces at the beginning or end, thereby ensuring consistency in text formatting and avoiding problems that can occur when searching or processing data. Example:

```
In [35]: import pandas as pd
data = {'Fruits': [' Apple ', 'Banana ', ' Cherry ', 'Date ']}
print(data)
{'Fruits': [' Apple ', 'Banana ', ' Cherry ', 'Date ']}
```

```
In [36]: df = pd.DataFrame(data)
```

```
In [37]: df['Fruits'] = df['Fruits'].str.strip()
```

```
In [38]: df.head()
```

```
Out[38]:
```

	Fruits
0	Apple
1	Banana
2	Cherry
3	Date

- **Inconsistent formatting:** Standardizing data formats, such as date formats, to make them consistent. Example, change the data format in the date column to YYYY-MM-DD for consistency

```
In [47]: df['Transaction_date'] = pd.to_datetime(df['Transaction_date'])

In [48]: df['Transaction_date'] = df['Transaction_date'].dt.strftime('%Y-%m-%d')

In [49]: df.head(10)

Out[49]:
```

	Transaction_date	Transaction_ID	Gender	Age	Marital_status	State_names	Segment	Employees_status	Payment_method	Referral	Amount_spent
0	2019-01-01	151200	Female	19.0	Single	Kansas	Basic	Unemployment	Other	Use	2051.360
1	2019-01-01	151201	Male	49.0	Single	Illinois	Basic	self-employed	Card	Not use	544.040
2	2019-01-01	151202	Male	63.0	Married	New Mexico	Basic	workers	PayPal	Use	1572.600
3	2019-01-01	151203	Female	18.0	Single	Virginia	Platinum	workers	Card	Use	1199.790
4	2019-01-01	151204	Male	27.0	Single	Connecticut	Basic	self-employed	Card	Not use	1341.435
5	2019-01-03	151205	Male	71.0	Single	Hawaii	Basic	Employees	PayPal	Use	2622.660
6	2019-01-03	151206	Female	34.0	Married	New Mexico	Platinum	Employees	PayPal	Use	1481.420
7	2019-01-03	151207	Male	37.0	Married	Connecticut	Basic	workers	PayPal	Use	1149.550
8	2019-01-04	151208	Male	75.0	Married	Florida	Silver	Employees	Card	Not use	1046.200
9	2019-01-04	151209	Female	41.0	Married	Vermont	Gold	Unemployment	Card	Use	2730.600

- **Removing duplicates:** Identifying and removing duplicate records to ensure data accuracy.

```
df.duplicated().sum()
df_new = df.drop_duplicates()
df_new.duplicated().sum()
```

Check for duplicate data in the table

```
In [13]: df.duplicated().sum()
```

```
Out[13]: 12
```

Removing the duplicates

```
In [14]: df_new = df.drop_duplicates()
```

Check data in the table after duplicate data has been deleted

```
In [15]: df_new.duplicated().sum()
```

```
Out[15]: 0
```

- **Imputing data:** Filling in missing data with appropriate values based on rules or algorithms.

```
df.Gender[df.Gender.isnull()]
df.Gender.value_counts()
val = df.Gender.mode().values[0]
df['Gender'] = df.Gender.fillna(val)
df.Gender.value_counts()
```

Check the amount of data/values in the categories in the Gender column

```
In [7]: df.Gender.value_counts()
```

```
Out[7]: Female    1356
      Male      1128
      Name: Gender, dtype: int64
```

From the proportion of the Gender column, Female is the data that appears most often, so Female is the mode

```
In [8]: val = df.Gender.mode().values[0]
      df['Gender'] = df.Gender.fillna(val)
```

After imputation, it can be seen that the proportions have changed

```
In [9]: df.Gender.value_counts()
```

```
Out[9]: Female    1384
      Male      1128
      Name: Gender, dtype: int64
```

- **Validating data:** Checking data for correctness and validity to ensure it meets predefined criteria or constraints

Before starting validation, it is important to examine the data by understanding the structure, data types, and any potential problems. It can use the head(), info(), and describe() methods of pandas DataFrame for this purpose.

df.head()

df.info()

In [62]: df.head()

Out[62]:

	Transaction_date	Transaction_ID	Gender	Age	Marital_status	State_names	Segment	Employees_status	Payment_method	Referral	Amount_spent
0	1/1/2019	151200	Female	19.0	Single	Kansas	Basic	Unemployment	Other	1.0	2051.360
1	1/1/2019	151201	Male	49.0	Single	Illinois	Basic	self-employed	Card	0.0	544.040
2	1/1/2019	151202	Male	63.0	Married	New Mexico	Basic	workers	PayPal	1.0	1572.600
3	1/1/2019	151203	Female	18.0	Single	Virginia	Platinum	workers	Card	1.0	1199.790
4	1/1/2019	151204	Male	27.0	Single	Connecticut	Basic	self-employed	Card	0.0	1341.435

In [28]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2512 entries, 0 to 2511
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Transaction_date       2512 non-null  object
1   Transaction_ID         2512 non-null  int64
2   Gender                 2512 non-null  object
3   Age                    2512 non-null  float64
4   Marital_status         2512 non-null  object
5   State_names            2512 non-null  object
6   Segment                2512 non-null  object
7   Employees_status       2512 non-null  object
8   Payment_method         2512 non-null  object
9   Referral               2512 non-null  float64
10  Amount_spent           2512 non-null  float64
dtypes: float64(3), int64(1), object(7)
memory usage: 216.0+ KB
```

Based on the results of the data inspection, it was found that the referral column used the float data type, this was deemed unsuitable, the data would be easier to understand if converted into a string, where "1.0" means using a referral code, while "0.0" means not using a code referral. This will help people who read the data so that it is easier to understand. (Data Type Validating)

df['Referral'] = df['Referral'].astype(str)

df['Referral'] = df['Referral'].replace({1.0: 'Use', 0.0: 'Not use'})

Check the data in the table after validating the data type

In [42]: df.head()

Out[42]:

	Transaction_date	Transaction_ID	Gender	Age	Marital_status	State_names	Segment	Employees_status	Payment_method	Referral	Amount_spent
0	1/1/2019	151200	Female	19.0	Single	Kansas	Basic	Unemployment	Other	Use	2051.360
1	1/1/2019	151201	Male	49.0	Single	Illinois	Basic	self-employed	Card	Not use	544.040
2	1/1/2019	151202	Male	63.0	Married	New Mexico	Basic	workers	PayPal	Use	1572.600
3	1/1/2019	151203	Female	18.0	Single	Virginia	Platinum	workers	Card	Use	1199.790
4	1/1/2019	151204	Male	27.0	Single	Connecticut	Basic	self-employed	Card	Not use	1341.435

In [43]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2512 entries, 0 to 2511
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Transaction_date       2512 non-null  object
1   Transaction_ID         2512 non-null  int64
2   Gender                 2512 non-null  object
3   Age                    2512 non-null  float64
4   Marital_status         2512 non-null  object
5   State_names            2512 non-null  object
6   Segment                2512 non-null  object
7   Employees_status       2512 non-null  object
8   Payment_method         2512 non-null  object
9   Referral               2512 non-null  object
10  Amount_spent           2512 non-null  float64
dtypes: float64(2), int64(1), object(8)
memory usage: 216.0+ KB
```

2.3 Organize data:

Purpose and common practices:

- **Sorting:** Reordering data based on one or more columns, usually in ascending or descending order.

example:

Sample data

```
data = {'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],  
       'Age': [25, 30, 22, 35, 28],  
       'Salary': [50000, 60000, 45000, 70000, 55000]}
```

Sorting: Reorder data based on the 'Age' column in ascending order

```
sorted_data = sorted(zip(data['Name'], data['Age'], data['Salary']), key=lambda x: x[1])  
print("Sorted data by Age:", sorted_data)
```

```
In [1]: # Sample data  
data = {'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],  
       'Age': [25, 30, 22, 35, 28],  
       'Salary': [50000, 60000, 45000, 70000, 55000]}
```

```
In [2]: # Sorting: Reorder data based on the 'Age' column in ascending order  
sorted_data = sorted(zip(data['Name'], data['Age'], data['Salary']), key=lambda x: x[1])  
print("Sorted data by Age:", sorted_data)
```

Sorted data by Age: [('Charlie', 22, 45000), ('Alice', 25, 50000), ('Eve', 28, 55000), ('Bob', 30, 60000), ('David', 35, 70000)]

- **Filtering:** Selecting a subset of data based on specified criteria.

example:

Filtering: Selecting records with Age greater than 25

```
filtered_data = [(name, age, salary) for name, age, salary in zip(data['Name'],  
data['Age'], data['Salary']) if age > 25]  
print("Filtered data:", filtered_data)
```

```
In [3]: # Filtering: Selecting records with Age greater than 25  
filtered_data = [(name, age, salary) for name, age, salary in zip(data['Name'], data['Age'], data['Salary']) if age > 25]  
print("Filtered data:", filtered_data)
```

Filtered data: [('Bob', 30, 60000), ('David', 35, 70000), ('Eve', 28, 55000)]

- **Slicing:** Extracting a specific range or portion of the data.

example:

Slicing: Extracting the second and third records

```
sliced_data = (data['Name'][1:3], data['Age'][1:3], data['Salary'][1:3])  
print("Sliced data:", sliced_data)
```

```
In [4]: # Slicing: Extracting the second and third records  
sliced_data = (data['Name'][1:3], data['Age'][1:3], data['Salary'][1:3])  
print("Sliced data:", sliced_data)
```

Sliced data: ('Bob', 'Charlie', [30, 22], [60000, 45000])

- **Transposing:** Changing the orientation of data, such as converting rows to columns or vice versa.

example:

Transposing: Changing rows to columns using zip

```
transposed_data = {'Name': data['Name'], 'Age': data['Age'], 'Salary': data['Salary']}  
print("Transposed data:", list(zip(*transposed_data.values())))
```

```
In [5]: # Transposing: Changing rows to columns using zip
transposed_data = {'Name': data['Name'], 'Age': data['Age'], 'Salary': data['Salary']}
print("Transposed data:", list(zip(*transposed_data.values()))))

Transposed data: [('Alice', 25, 50000), ('Bob', 30, 60000), ('Charlie', 22, 45000), ('David', 35, 70000), ('Eve', 28, 55000)]
```

- **Appending:** Combining or adding new data to an existing dataset.

example:

```
# Appending: Adding new data to the existing dataset
new_data = {'Name': ['Frank', 'Grace'], 'Age': [29, 32], 'Salary': [52000, 60000]}
appended_data = {key: data[key] + new_data[key] for key in data}
print("Appended data:", appended_data)
```

```
In [6]: # Appending: Adding new data to the existing dataset
new_data = {'Name': ['Frank', 'Grace'], 'Age': [29, 32], 'Salary': [52000, 60000]}
appended_data = {key: data[key] + new_data[key] for key in data}
print("Appended data:", appended_data)

Appended data: {'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve', 'Frank', 'Grace'], 'Age': [25, 30, 22, 35, 28, 29, 32], 'Salary': [50000, 60000, 45000, 70000, 55000, 52000, 60000]}
```

- **Truncating:** Reducing the data to a specific length or number of rows, often to create smaller subsets of data.

example:

```
# Truncating: Reducing the dataset to the first three records
truncated_data = {key: data[key][:3] for key in data}
print("Truncated data:", truncated_data)
```

```
In [7]: # Truncating: Reducing the dataset to the first three records
truncated_data = {key: data[key][:3] for key in data}
print("Truncated data:", truncated_data)

Truncated data: {'Name': ['Alice', 'Bob', 'Charlie'], 'Age': [25, 30, 22], 'Salary': [50000, 60000, 45000]}
```

2.4 Aggregate data:

Purpose and common practices:

- **Grouping:** Grouping data by one or more columns to perform operations on subsets of data, often used with aggregation functions like SUM, COUNT, AVG.

example:

1. SUM

```
import pandas as pd
```

```
data = {
    'Category': ['Electronics', 'Clothing', 'Electronics', 'Clothing', 'Electronics'],
    'Amount': [500, 200, 300, 150, 700]
}
```

```
df = pd.DataFrame(data)
```

```
# Grouping by 'Category' and calculating the total sales (SUM)
grouped_data = df.groupby('Category')['Amount'].sum()
print(grouped_data)
```

```
In [8]: import pandas as pd

data = {
    'Category': ['Electronics', 'Clothing', 'Electronics', 'Clothing', 'Electronics'],
    'Amount': [500, 200, 300, 150, 700]
}

df = pd.DataFrame(data)

# Grouping by 'Category' and calculating the total sales (SUM)
grouped_data = df.groupby('Category')['Amount'].sum()
print(grouped_data)

Category
Clothing      350
Electronics  1500
Name: Amount, dtype: int64
```

2. COUNT

```
count_data = df.groupby('Category')['Amount'].count()
print(count_data)
```

```
In [9]: count_data = df.groupby('Category')['Amount'].count()
print(count_data)

Category
Clothing      2
Electronics   3
Name: Amount, dtype: int64
```

3. AVG

```
average_data = df.groupby('Category')['Amount'].mean()
print(average_data)
```

```
In [10]: average_data = df.groupby('Category')['Amount'].mean()
print(average_data)

Category
Clothing      175.0
Electronics   500.0
Name: Amount, dtype: float64
```

- **Joining/Merging:** Combining data from multiple sources or tables using keys or common columns.

example:

```
customer_data = {
    'CustomerID': [1, 2, 3],
    'Name': ['Alice', 'Bob', 'Charlie']
}
```

```
purchase_data = {
    'CustomerID': [2, 1, 3],
    'Product': ['Laptop', 'Phone', 'Tablet']
}
```

```
customers = pd.DataFrame(customer_data)
purchases = pd.DataFrame(purchase_data)
```

```
# Merge the two DataFrames on 'CustomerID'
merged_data = pd.merge(customers, purchases, on='CustomerID')
print(merged_data)
```



```
In [11]: customer_data = {
        'CustomerID': [1, 2, 3],
        'Name': ['Alice', 'Bob', 'Charlie']
    }

    purchase_data = {
        'CustomerID': [2, 1, 3],
        'Product': ['Laptop', 'Phone', 'Tablet']
    }

    customers = pd.DataFrame(customer_data)
    purchases = pd.DataFrame(purchase_data)

    # Merge the two DataFrames on 'CustomerID'
    merged_data = pd.merge(customers, purchases, on='CustomerID')
    print(merged_data)
```

	CustomerID	Name	Product
0	1	Alice	Phone
1	2	Bob	Laptop
2	3	Charlie	Tablet

- **Summarizing:** Creating summary statistics or aggregations to get an overview of the data, such as calculating totals, averages, or counts.

example:

```
summary_stats = df.describe()
print(summary_stats)
```

```
In [12]: summary_stats = df.describe()
        print(summary_stats)
```

	Amount
count	5.000000
mean	370.000000
std	228.035085
min	150.000000
25%	200.000000
50%	300.000000
75%	500.000000
max	700.000000

- **Pivoting:** Restructuring data to transform rows into columns or vice versa, often used for creating summary tables or pivot tables.

example:

```
pivoted_data = df.pivot(index='Category', columns='Amount', values='Amount')
print(pivoted_data)
```

```
In [13]: pivoted_data = df.pivot(index='Category', columns='Amount', values='Amount')
        print(pivoted_data)
```

Amount	150	200	300	500	700
Category					
Clothing	150.0	200.0	NaN	NaN	NaN
Electronics	NaN	NaN	300.0	500.0	700.0